

Girish Bharambe
NVIDIA

PROGRAM EXECUTION ENVIRONMENT



WHY LEARN PROGRAM EXECUTION ENVIRONMENT

- Learn details of what is going on “under the hood” of a computer system
 - Being able to understand various abstractions and interfaces that exist between various system software components and Hardware
 - Not just “what” but also “why” and “how”
 - Being able to understand how does computer execute your programs
- Why should you learn these topics?
 - Makes you a better programmer
 - Makes you more effective at debugging
 - Makes you aware of how system is built and that enables you to build new systems as needed
 - Note although specific implementations of systems change, underlying concepts mostly hold or at least form initial intuitions for new ideas

Life of a program

Abstraction of functions

Mechanisms to implement Functions

Mechanisms for Control Transfer


Mechanism for Local Variables

Mechanism for Parameters

Mechanism for Register Usage

Register Saving Conventions

OUTLINE



LIFE OF A PROGRAM



LIFE OF A PROGRAM (HIGH LEVEL VIEW)


```
int foo(int p) {  
    int var = 5;  
    return var + p;  
}
```



```
foo:  
    pushl %ebp  
    movl %esp, %ebp  
    ...  
    addl %edx, %eax  
    ret
```



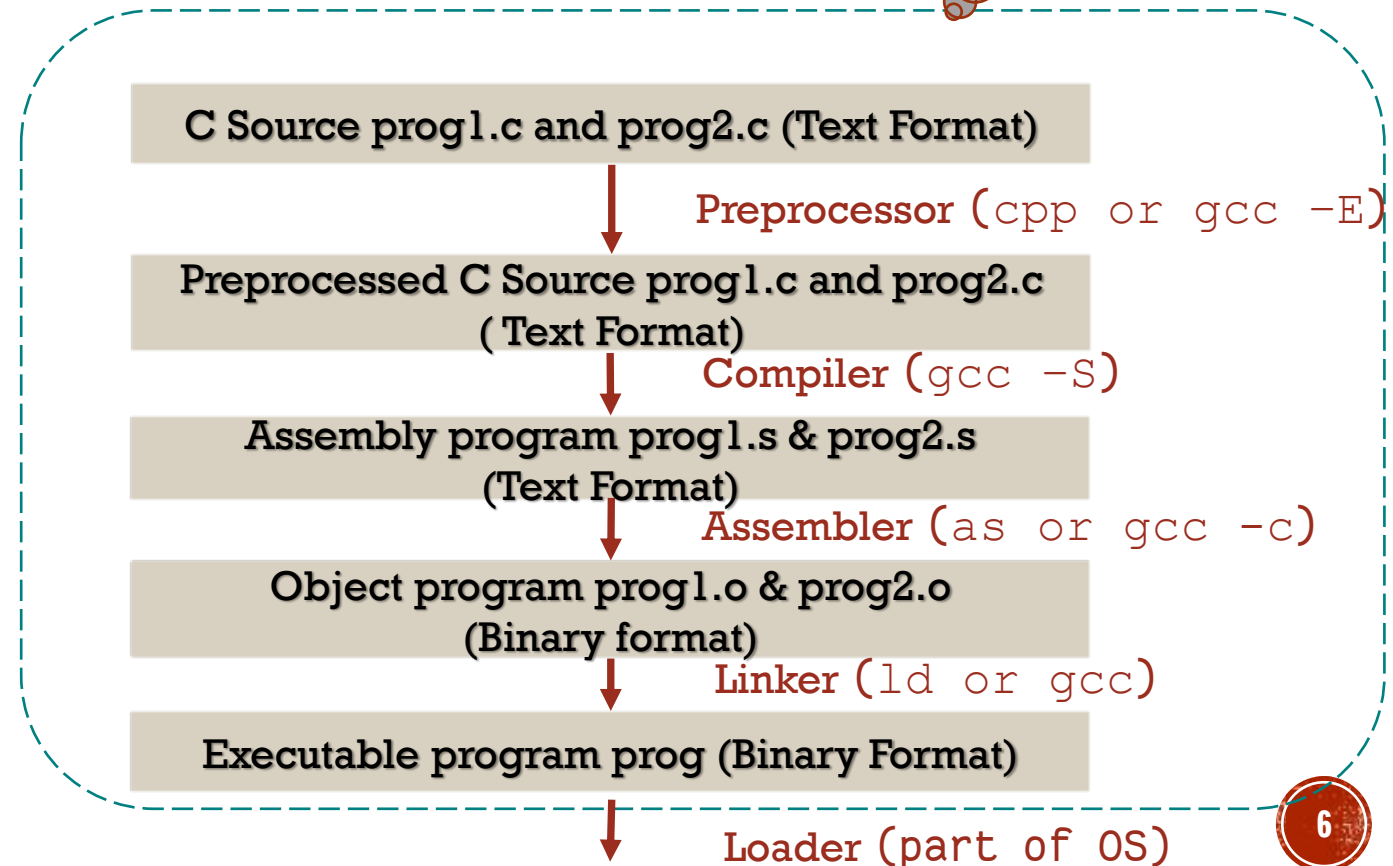
```
1010101  
1000100111100101  
...  
000111010000  
11000011
```



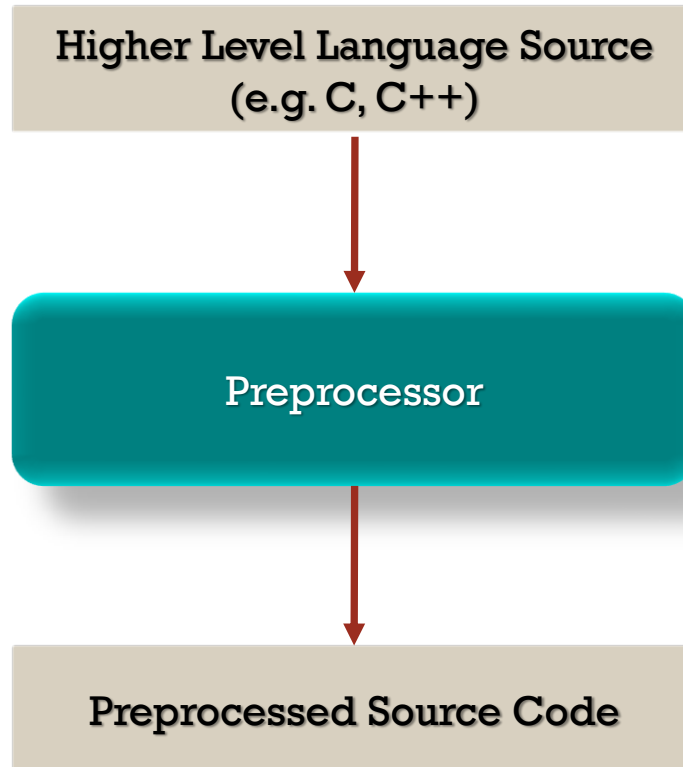
BUILDING AN EXECUTABLE FROM C SOURCE

Compiler
Driver

- Consider source code in prog1.c and prog2.c
- Compile the source file into executable using
 - `gcc prog1.c prog2.c -o prog`
- Run the program using
 - `./prog`

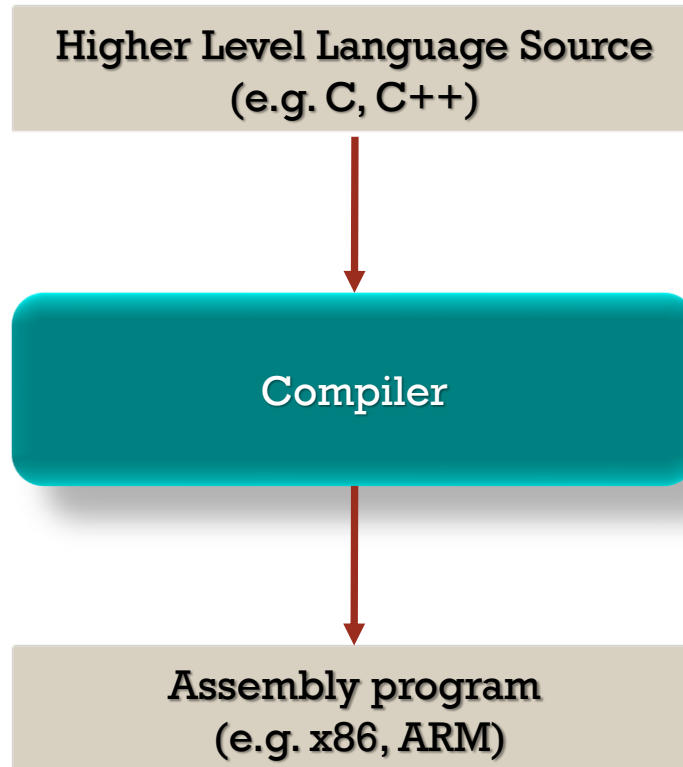


PREPROCESSOR




- Processes pre-processor directives like `#define`, `#include`, macro expansions.
- Typically is a separate program invoked before actual compilation.

COMPILER




- Translates source written in higher level languages into lower level assembly
- Not just translator but also has optimizations and code generation

ASSEMBLER



- Translate textual assembly program into binary object file
- Assembly languages typically support various directives to write different parts of assembly program e.g. `.text`, `.data` which needs to be handled by assemblers
- Encodes textual machine instructions to corresponding instruction encoding
- Creates object files containing machine instructions and other required information

LINKER

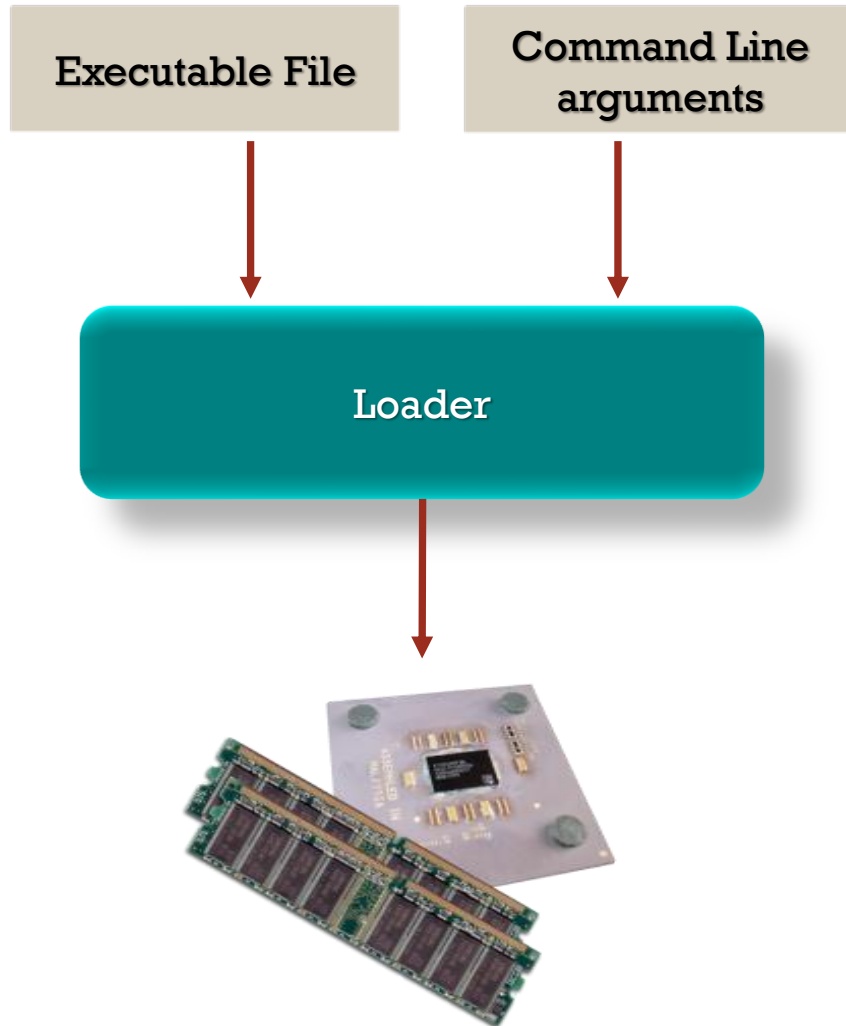


- Combines several object files into a single executable
- Primarily need to resolve extern references from object files to their corresponding definitions from other object file.
- Creates executable binary program that can be loaded by loader.

WHAT IS NEED FOR A LINKER?

- Recall : Linker combines several object files into a single executable
- **Why write source code in multiple files?**
 - Very uncommon to have large codes to be fully self-contained without any other dependencies.
 - For modularity, better to write large program as collection of smaller source files
- Code Reuse
 - Allows reuse of already written codes
 - Can create library of common utilities
 - Facilitates creation and distribution of pre-compiled libraries
 - IP protection : Allows library author to distribute code without shipping source code
- Separate Compilation
 - When one of the source file changes, no need to recompile all parts of program, can only compile modified source file and relink
- **BUT WAIT, are there NO advantages of having source code in single file?**
 - Having source in single file provides more visibility to compiler which enables more optimizations
 - Linker will have visibility and can do some optimizations at link time but is limited compared to compiler optimizations

LOADER

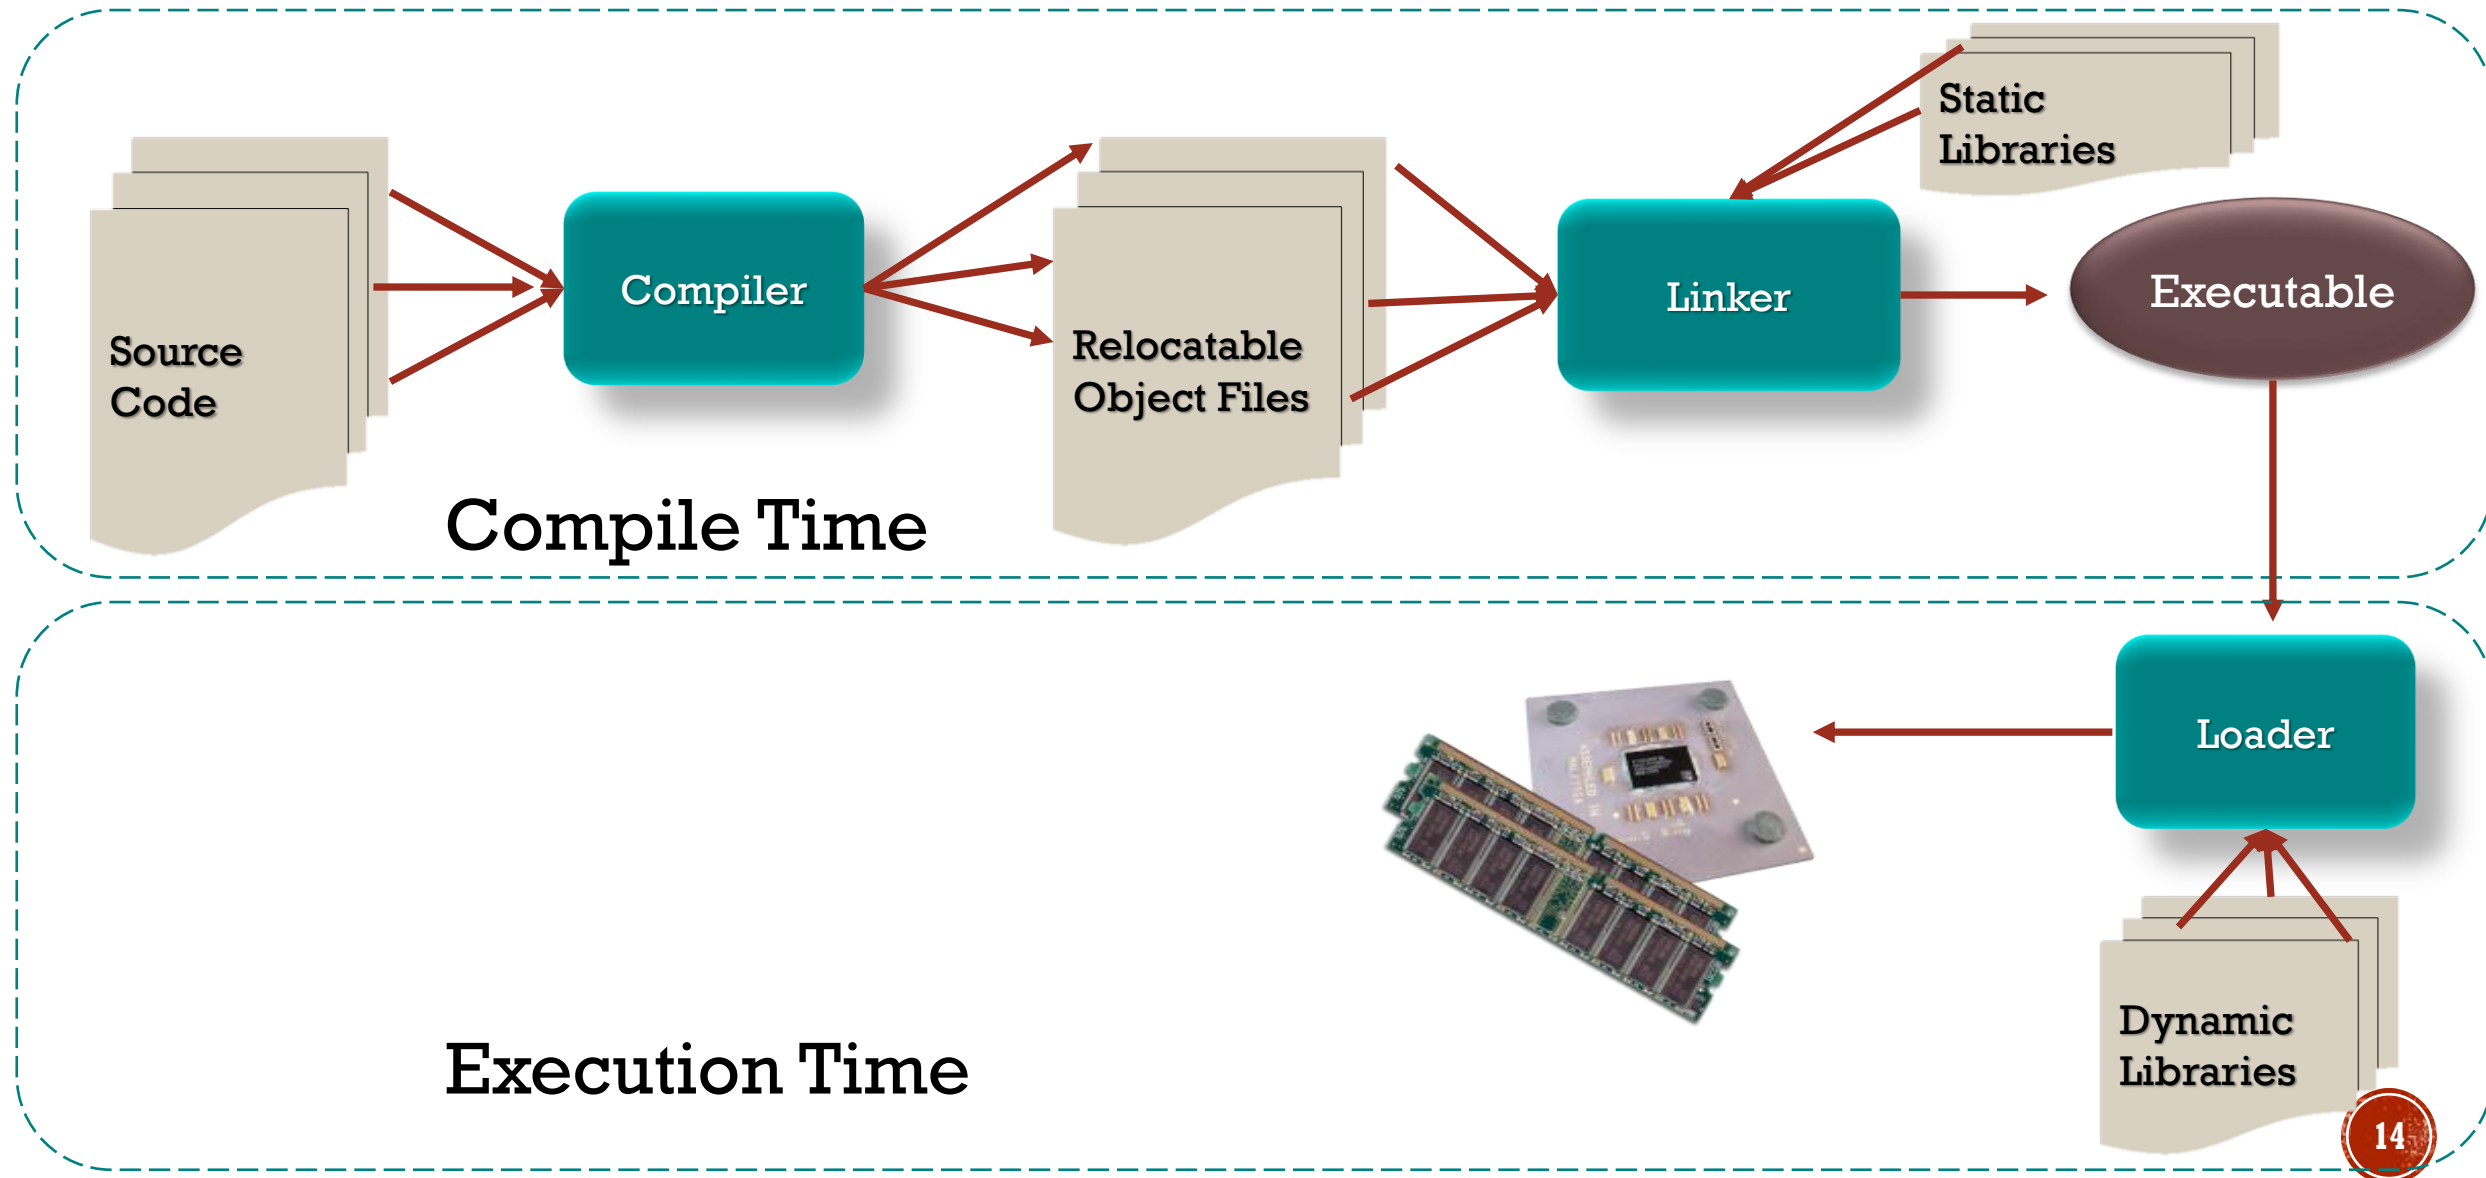



- Part of OS / Kernel
- Reads executable binary from disk and loads it into memory and then starts running program by jumping to first instruction of program

WHAT IS NEED FOR A LOADER?

- In a system that supports running only one program at a time, may be no need for a loader
 - Program can be assembled and linked for a fixed memory address
 - Program has access to entire memory
- With operating system and system running multiple programs, physical memory is shared between various programs including operating systems
- Actual address at which program will be loaded isn't known until execution time

LIFE OF A PROGRAM





DEMYSTIFYING ABSTRACTION OF FUNCTIONS

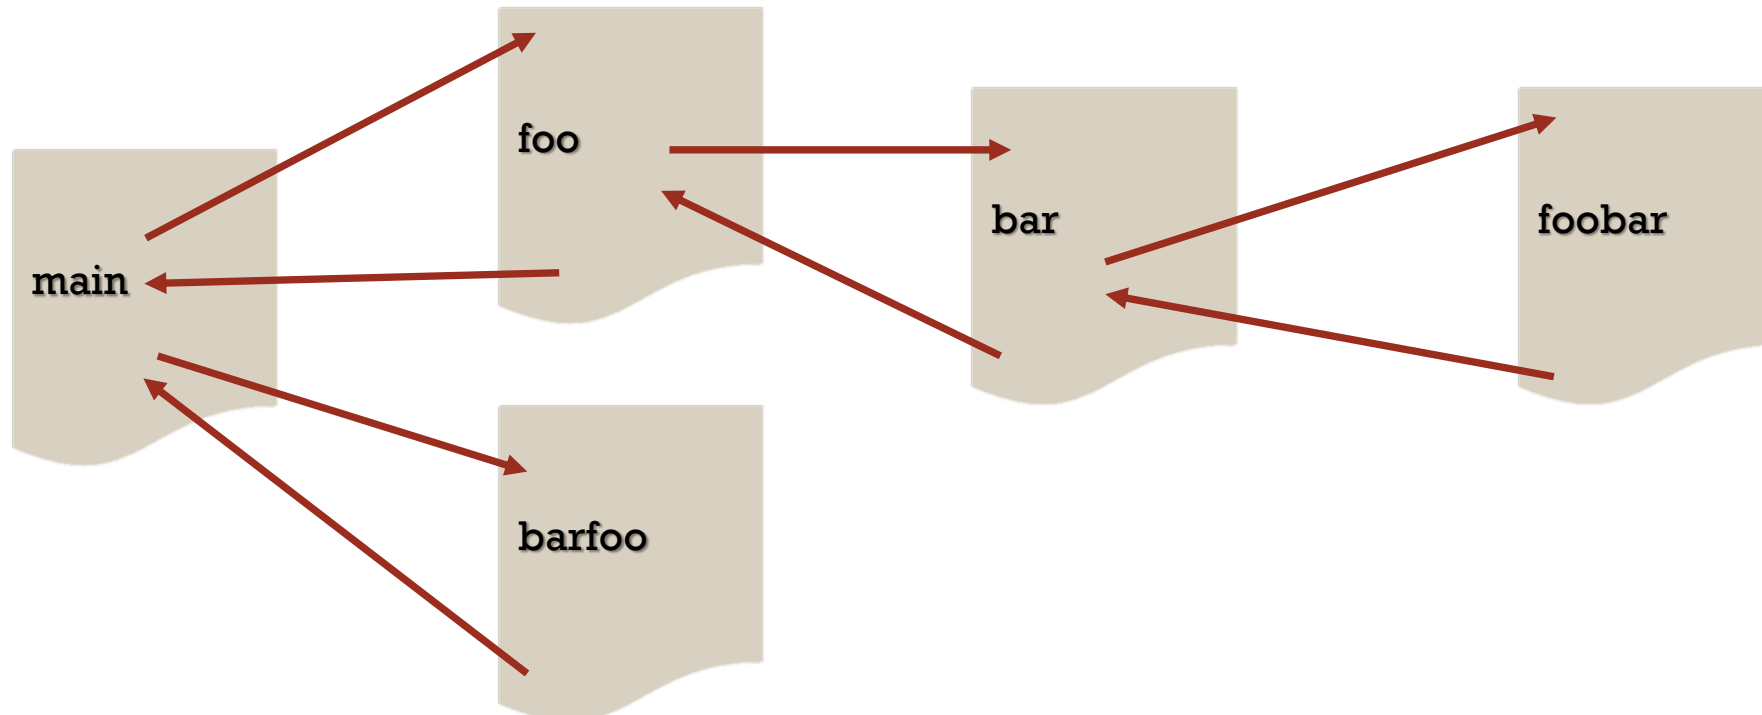


FUNCTIONS

- Programming Languages allow encapsulating a block of code that *typically* perform single, related task.
- Functions allow better modularity and provide high degree of code reuse
- A function consists of a
 - **Name** that is symbol that represents address of function
 - **Parameters** (aka input arguments) that are used to specify data that is necessary for function to work
 - **Return value** (aka output argument) that is used to return data back to the caller function
 - **Function body** that contains collection of instructions that perform operations function is intended to do.
 - **Local Variables** that are data storage that function uses during processing and is thrown away when function returns.

CHARACTERISTICS OF FUNCTION

- Each function *typically* has a single-entry point
- The calling function (aka *caller*) *typically* is suspended during execution of the called function (aka *callee*)
- Control returns to the caller function when callee function returns i.e. function returns only after all functions it calls have returned

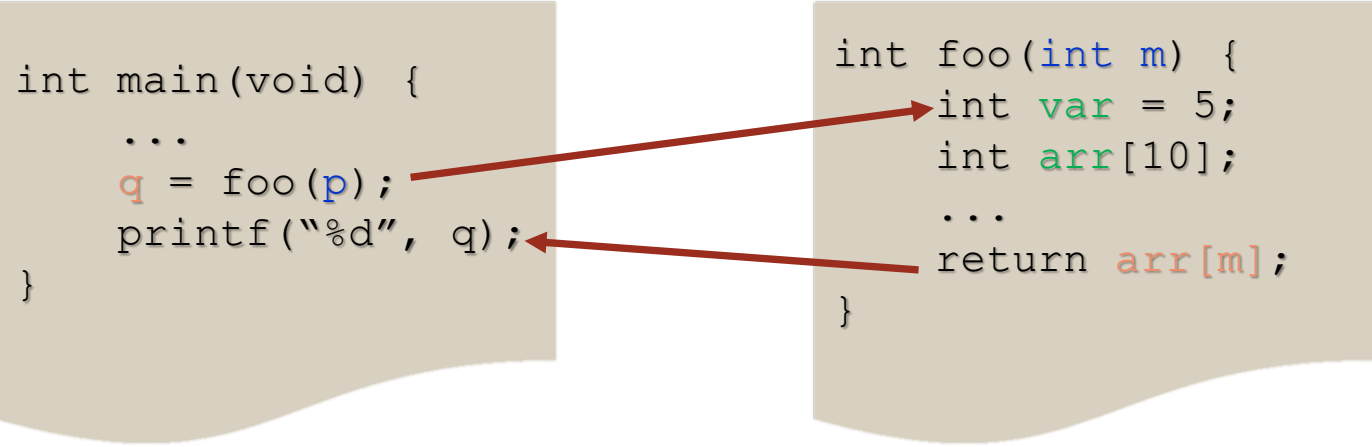


MECHANICS NEEDED TO IMPLEMENT FUNCTION CALL

- Mechanics to **transfer control** from caller function to callee and return from callee function to caller
- Mechanics to pass **function parameters** and **return values**
- Mechanics to allocate memory for **local variables** and deallocate after function returns
- Mechanics for register usage across caller and callee functions

```
int main(void) {  
    ...  
    q = foo(p);  
    printf("%d", q);  
}
```

```
int foo(int m) {  
    int var = 5;  
    int arr[10];  
    ...  
    return arr[m];  
}
```



MECHANICS FOR CONTROL TRANSFER

- Function call requires processor to start executing instructions from callee and when returned from callee, continue execution of instructions following call instruction
- This mechanism can be implemented using simple branch / jump instruction.
 - X86 processor has `jmp` instruction which jumps execution control to specified address

```
int main() {  
    foo(5);  
}  
int foo(int p) {  
    return p + 5;  
}
```



```
main:  
    ...  
    jmp foo;  
    ...  
foo:  
    ...
```


MECHANICS FOR CONTROL TRANSFER (CONT)

- But what about transferring control back to caller?
 - How does callee know where to return to?
- Save address of where to return to before doing the call
- Use this saved address to jump back
- But which address really?
- Where to save it?

RETURN ADDRESS AS IMMEDIATE/LABEL

- Do we really need to save address? Why not simply jmp back to required address?

```
int main() {  
    foo(5);  
}  
int foo(int p) {  
    return p + 5;  
}
```



- **What are issues with this scheme?**
 - jmp in foo will always return to same address.
 - Won't work if foo is called multiple times from main or from different function.
 - Return address is different for every invocation of a function

RETURN ADDRESS IN GLOBAL VARIABLE

- Say we create a dummy global variable `retAddr_foo` that holds return address

```
int main() {  
    foo(5);  
}  
int foo(int p) {  
    return p + 5;  
}
```



```
.bss  
retAddr_foo:  
    .space 4  
main:  
    movl lbl, retAddr_foo  
    jmp foo  
lbl:  
    ...  
foo:  
    ...  
    jmp *retAddr_foo
```


- **What are issues with this scheme?**
 - Both caller and callee function must agree on name and location of global variable. May not be best interface design.
 - **Will this work for all cases?**
 - How will recursive functions work in this model?

HANDLING OF RETURN ADDRESS (CONT)

- So far, we know:
 - Both caller and callee **must agree** on location of return address.
 - Storage used for return address is specific to every invocation of the function and hence must be created for every function invocation

STACK

- Many programming languages support stack-based execution for supporting recursion.
- Stack is a region of memory that works in LIFO manner.
- x86 and many other modern processors have native support for stack in HW.



X86 STACK


- HW register `%esp` acts as stack pointer
- Push new data on stack
 - Instruction `pushl src`
 - Store value at address given by `%esp`
 - Decrement `%esp` by 4
 - `pushl %eax` equivalent to

```
subl $4, %esp
movl %eax, (%esp)
```
- Pop data from stack
 - Instruction `popl dst`
 - Loads value at address given by `%esp` in `dst`
 - Increments `%esp` by 4
 - `popl %eax` Equivalent to

```
movl %(esp), %eax
addl $4, %esp
```


Higher Address

Lower Address



X86_64 STACK

- HW register `%rsp` acts as stack pointer
- Push new data on stack
 - Instruction `pushq src`
 - Store value at address given by `%rsp`
 - Decrement `%rsp` by 8
- Pop data from stack
 - Instruction `popq dst`
 - Loads value at address given by `%rsp` in `dst`
 - Increments `%rsp` by 8

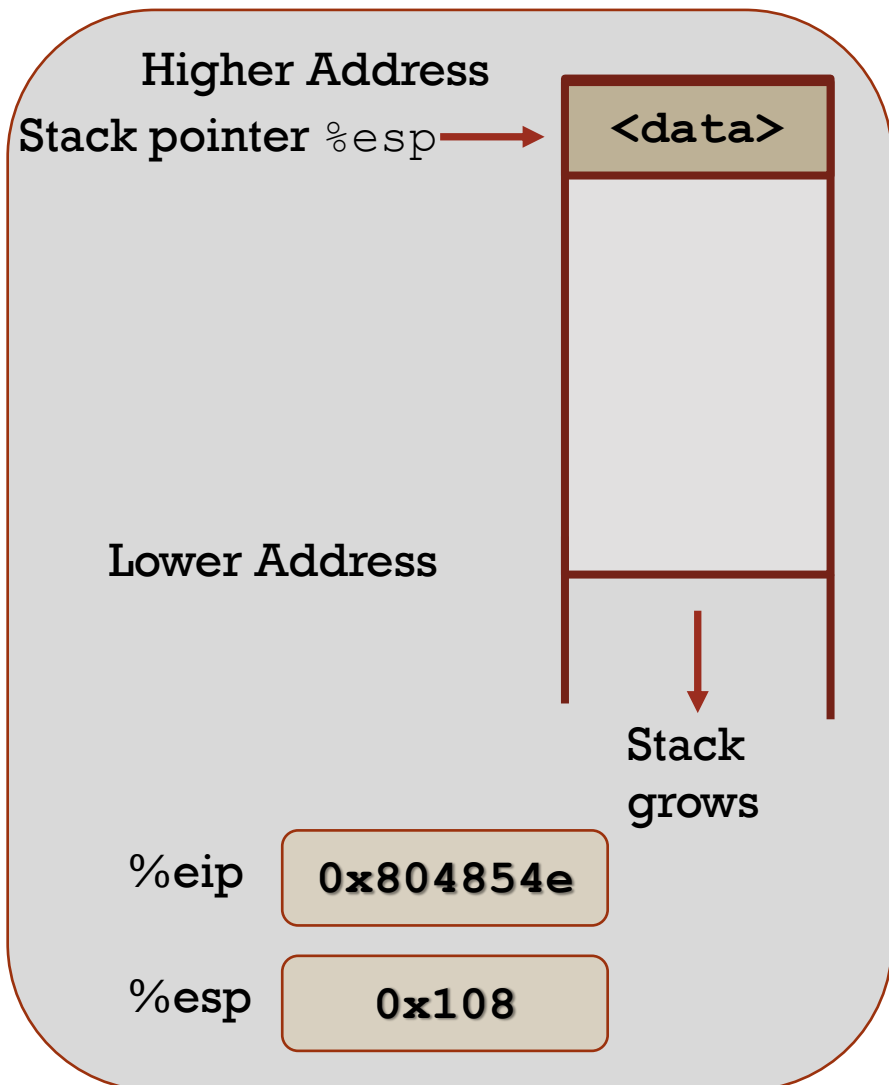


MECHANICS FOR CONTROL TRANSFER IN X86

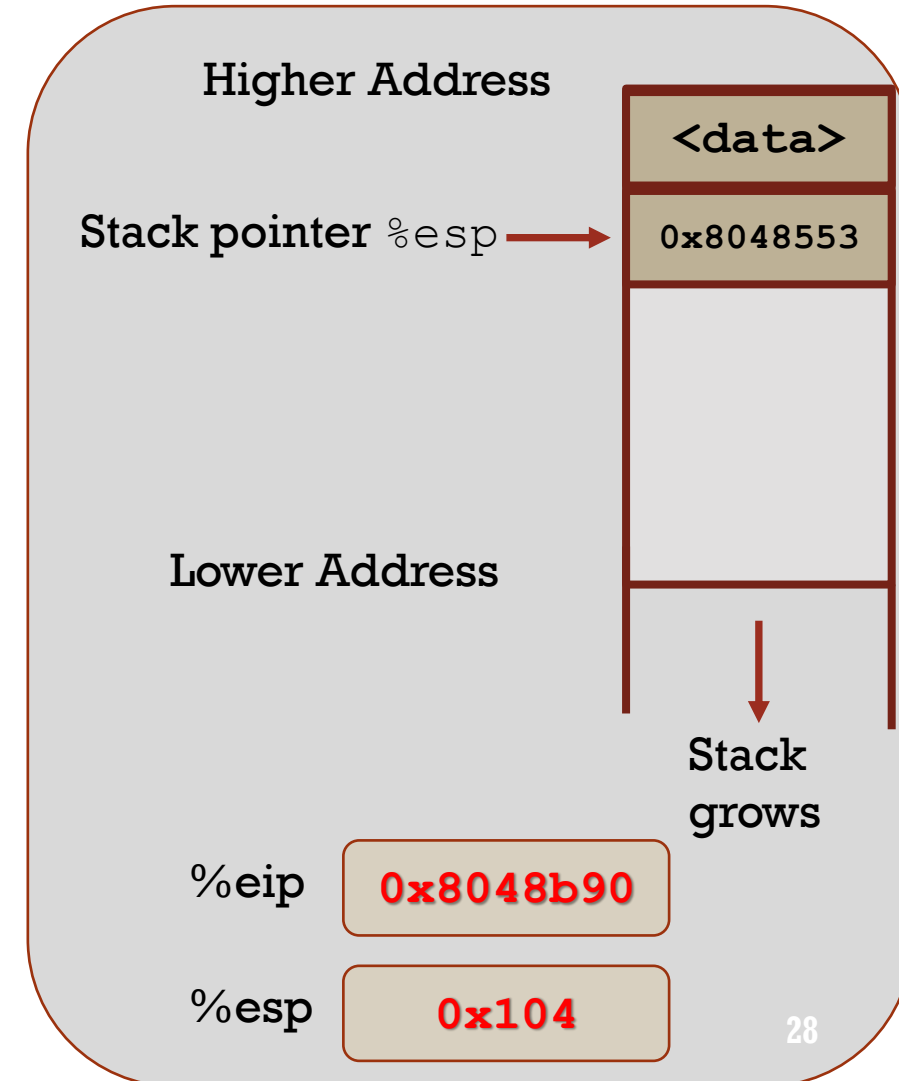
- x86 ISA supports following instructions for implementing function calls and returns

| Instruction | Description |
|-------------|---|
| call Label | Function call <ul style="list-style-type: none">• Pushes return address on the stack and jump to start of the called function |
| ret | Return from call <ul style="list-style-type: none">• Pops address from top of stack and jumps to that location• Proper use requires just before instruction, stack is setup such that stack pointer points to place where call stored return address |

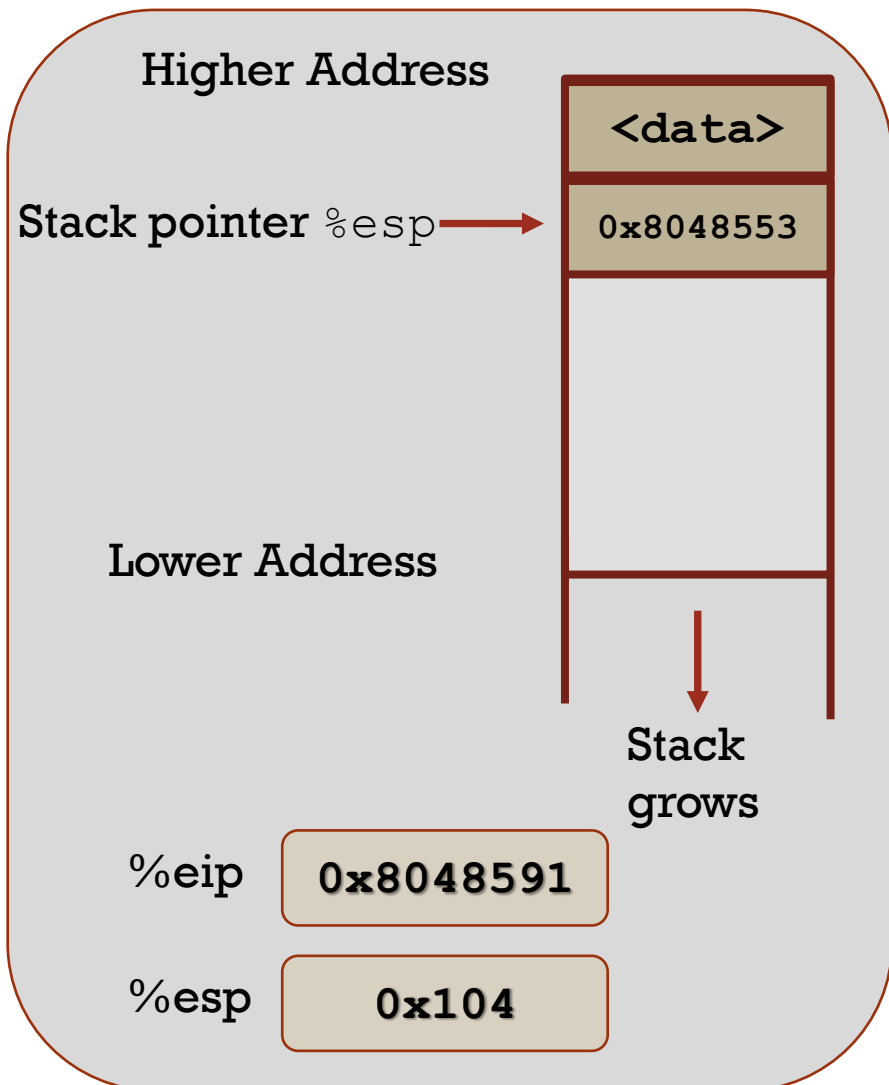
FUNCTION CALL DEMONSTRATION



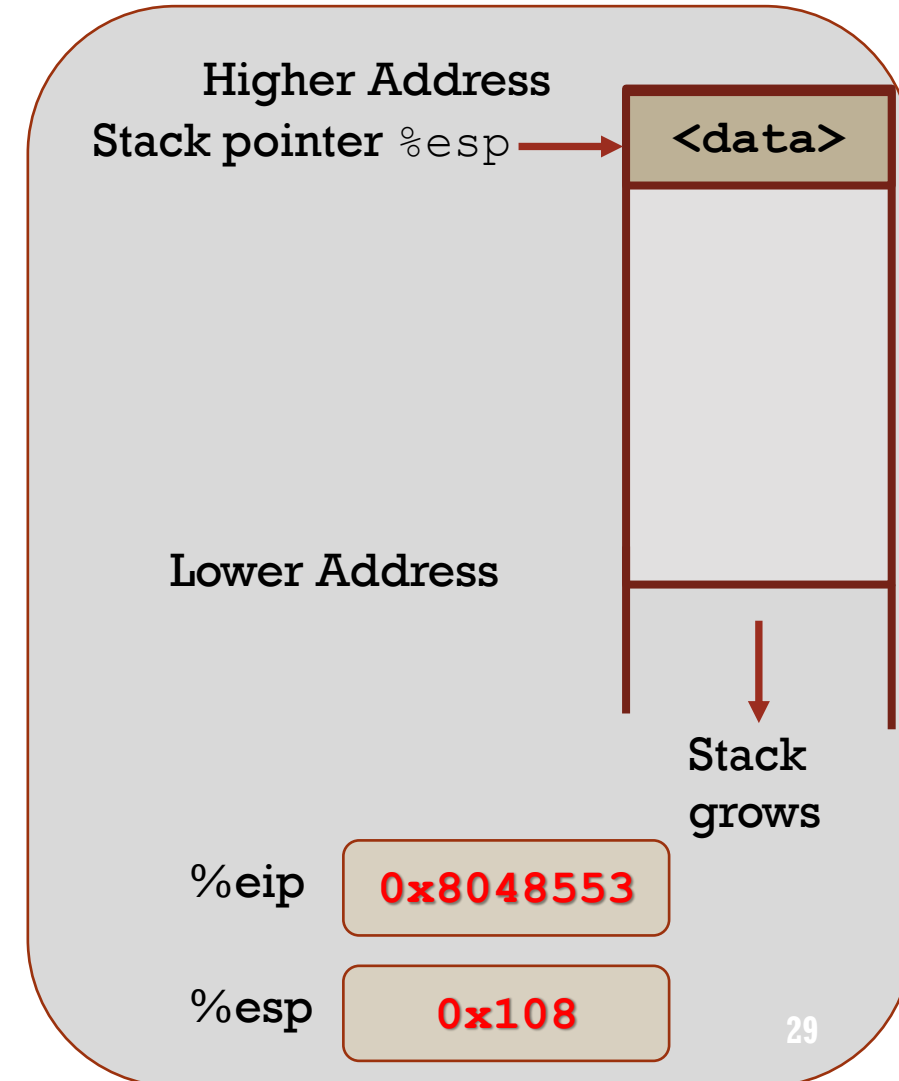
```
0x804854e: call 8048b90
0x8048553: movl $0, %eax
```



FUNCTION RETURN DEMONSTRATION



```
0x8048591: ret
```



MECHANICS FOR LOCAL VARIABLES IN FUNCTION

- A function can define its own variables often referred to as local variables
- Languages like C/C++ support two types of local variables
 - **Variables** with `auto` storage (default)
 - **Variables** with `static` storage
- What is the scope of function local variables?
 - Local variables are accessible only in the function
- What is the lifetime of function local variables?
 - Local variables with `auto` storage have the lifetime of the function
 - `static` local variables have the lifetime of the entire program
- When are `auto` variables allocated?
 - Each invocation of the function has its own instantiation of `auto` variables
 - Dynamically allocated, initialized every time the function is called
 - Deallocated every time the function returns
- When are `static` variables allocated?
 - When the program is loaded, details will not be covered in this talk

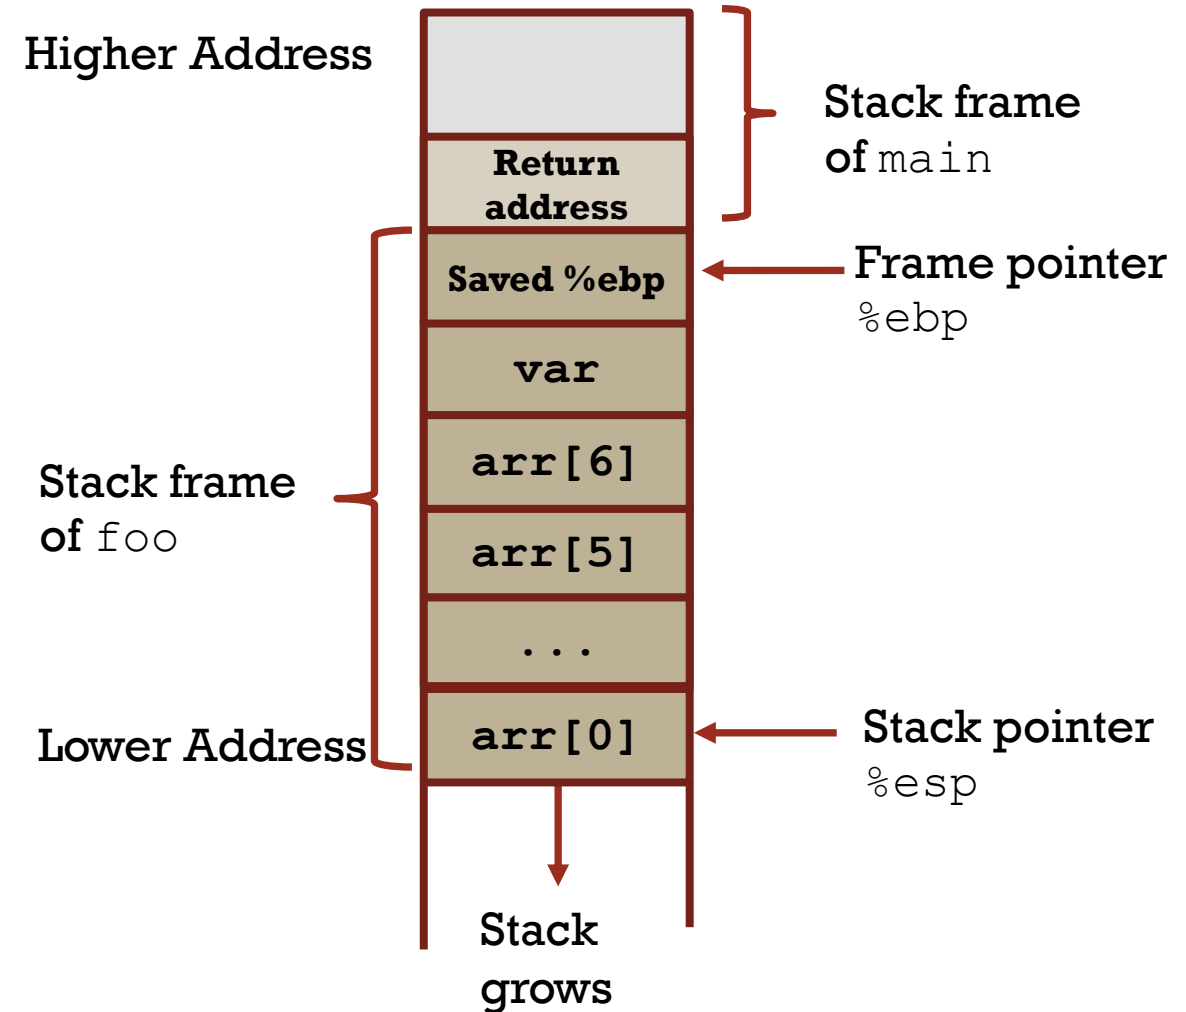
```
int foo(int m) {  
    int var = 5;  
    int arr[10];  
    static int cnt;  
    ...  
    return arr[m];  
}
```

STACK FRAME IN X86

- Portion of stack allocated for a function call is called “*stack frame*”.
 - An implementation of “*Activation Records*” for stack-based languages
- Apart from stack pointer, also has frame pointer `%ebp`

```
int main(void) {  
    ...  
    q = foo();  
    printf("%d", q);  
}
```

```
int foo() {  
    int var = 5;  
    int arr[7];  
    ...  
    return arr[var];  
}
```



ALLOCATE SPACE IN STACK

- Function `foo` requires $4 + 7 * 4 = 32$ bytes of stack for holding auto local variables.

```
foo:
    subl $32, %esp
    ...
```

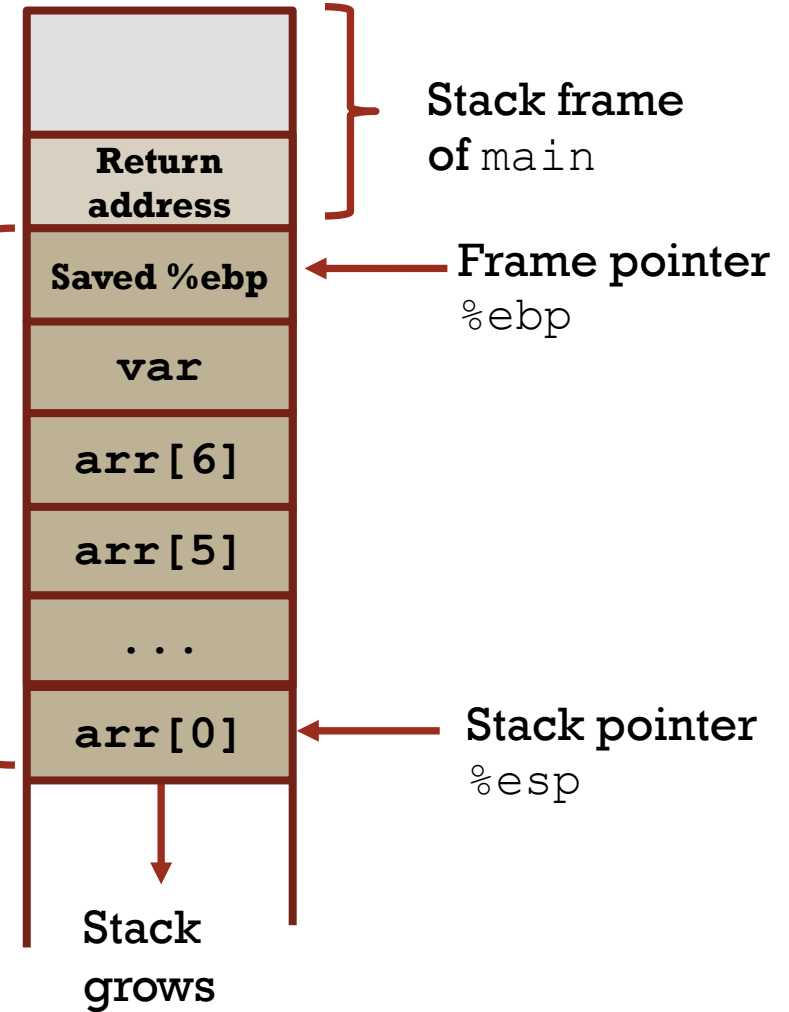
```
int main(void) {
    ...
    q = foo();
    printf("%d", q);
}
```

```
int foo() {
    int var = 5;
    int arr[7];
    ...
    return arr[var];
}
```

Higher Address

Stack frame
of `foo`

Lower Address



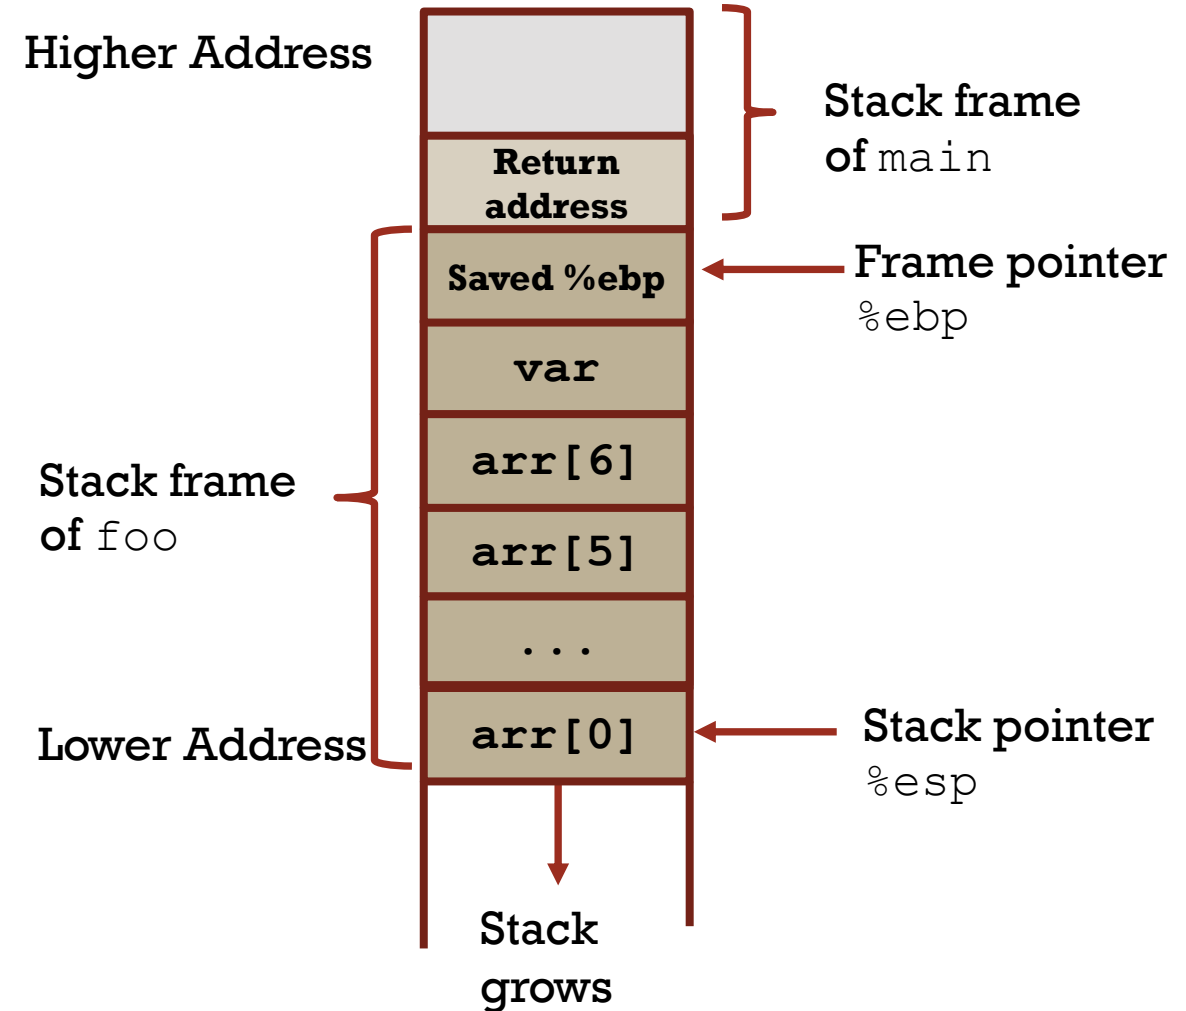
ACCESS AUTO LOCAL VARIABLES IN STACK

- How does program access `var`?
 - What is address of `var`?

```
foo:
    subl $32, %esp
    ...
    movl    $5, 28(%esp)
```

```
int main(void) {
    ...
    q = foo();
    printf("%d", q);
}
```

```
int foo() {
    int var = 5;
    int arr[7];
    ...
    return arr[var];
}
```

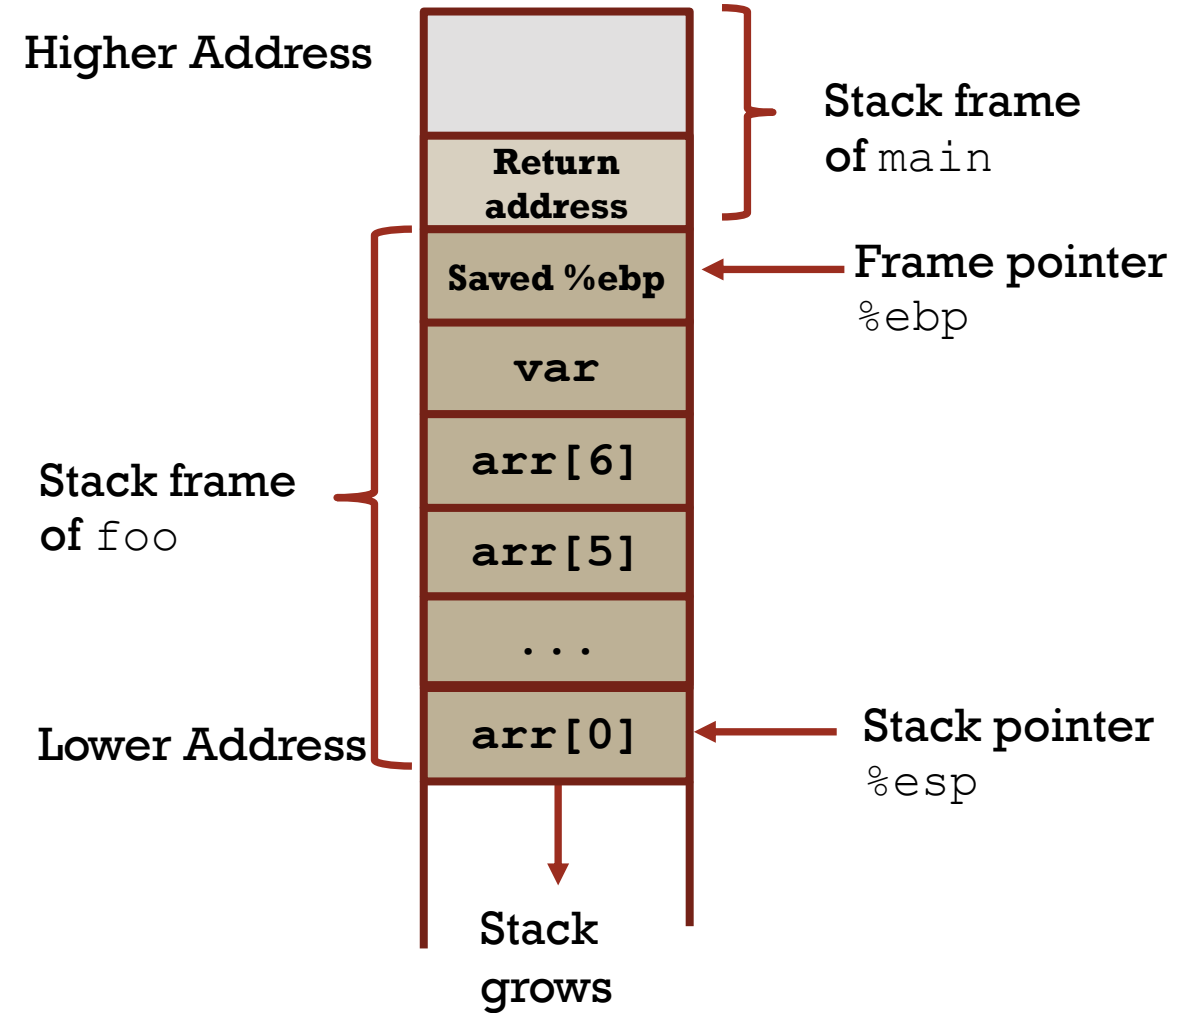


FRAME POINTER

- Why is frame pointer needed?
 - As stack grows, address of variables relative to stack pointer will change
 - Variables always at fixed address from frame pointer, `var` will always be at `%ebp - 4`
 - `var` can be accessed as
`movl $5, -4(%ebp)`

```
int main(void) {  
    ...  
    q = foo();  
    printf("%d", q);  
}
```

```
int foo() {  
    int var = 5;  
    int arr[7];  
    ...  
    return arr[var];  
}
```




PREPARE STACK FOR RETURN

- What all is needed to return to caller?
 - Deallocate stack allocated for local variables etc.
 - Need to ensure stack pointer is pointing to return address

```
foo:
    ...
    movl %ebp, %esp
    popl %ebp
    ret
```

```
int main(void) {
    ...
    q = foo();
    printf("%d", q);
}
```

```
int foo() {
    int var = 5;
    int arr[7];
    ...
    return arr[var];
}
```



FUNCTION PROLOGUE AND EPILOGUE

| Function Prologue | Function Epilogue |
|---|---|
| Lines of code at beginning of function to prepare stack and registers for use within the function | Reverses action of prologue and returns control to caller function |
| Instruction <code>enter N, 0</code> is equivalent to prologue shown below | Instruction <code>leave</code> is equivalent <code>movl + popl</code> instruction |

```
foo:
    pushl %ebp          ; save %ebp
    movl %esp, %ebp    ; update %ebp
    subl $<N>, %esp    ; space in stack
    ...
```

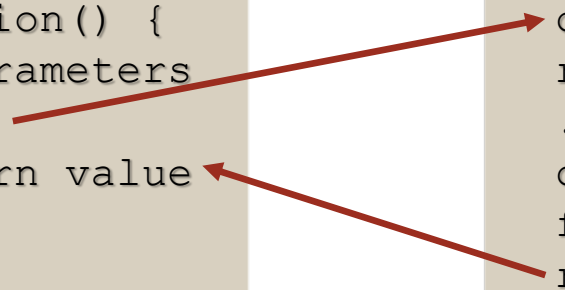
```
foo:
    ...
    movl %ebp, %esp    ; Update %esp
    popl %ebp          ; restore %ebp
    ret                ; return
```

MECHANICS FOR FUNCTION PARAMETERS AND RETURN VALUES

- How does callee know where to find parameters?
- How does caller know where to find return value?
- Need to have some contract between caller and callee function for this.

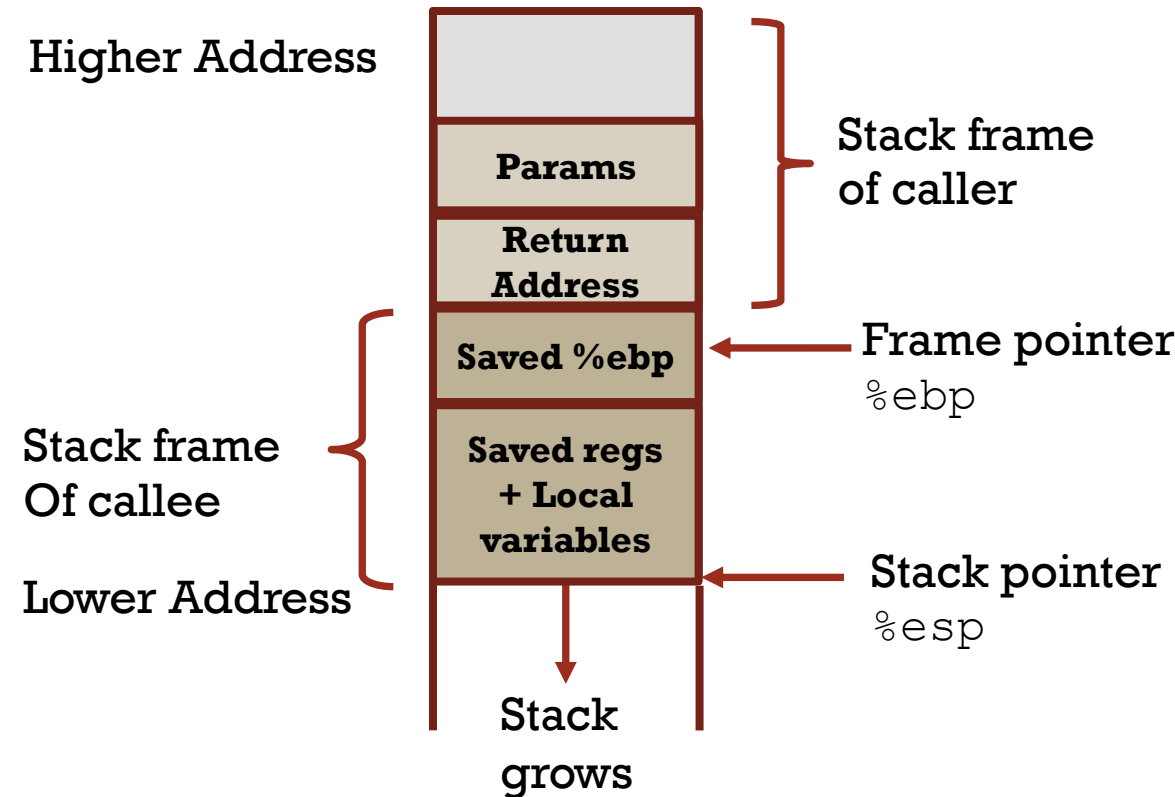
```
CallerFunction() {  
  set up parameters  
  call  
  read return value  
}
```

```
CalleeFunction() {  
  create local vars  
  read parameters  
  ...  
  create return value  
  free local vars  
  ret  
}
```



STACK FRAME X86 REVISITED

- Caller function allocates parameters in stack frame just before call instruction which will push return address on stack
- This forms **yet another contract** between caller and callee



STACK FRAME IN X86 (CONT)

- How to access parameter `m`?
 - `%ebp + 8`

```
int main(void) {  
    ...  
    r = foo(p, q);  
    printf("%d", r);  
}
```

```
int foo(int m, int n) {  
    int x = m + n;  
    int y = 5;  
    int sum = x+y;  
    ...  
    return sum;  
}
```

Higher Address

`%ebp+12`

`%ebp+8`

`%ebp+4`

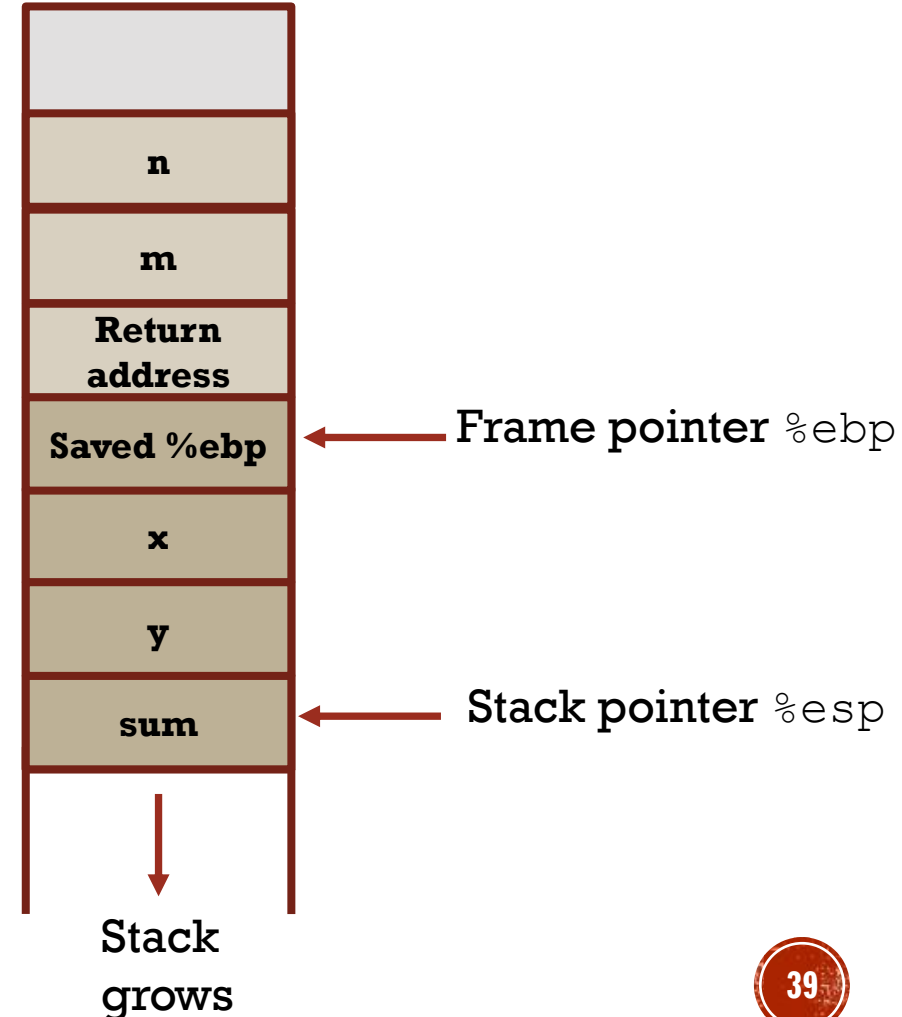
`%ebp`

`%ebp-4`

`%ebp-8`

`%ebp-12`

Lower Address



STACK FRAME IN X86 (CONT)

- What should be order of parameters in the stack?
 - Should they be passed from right to left or left to right or some other order?
 - Does this even matter as long as caller and callee agree on order?
 - Varargs functions

```
int main(void) {  
    ...  
    r = foo(p, q);  
    printf("%d", r);  
}
```

```
int foo(int m, int n) {  
    int x = m + n;  
    int y = 5;  
    int sum = x+y;  
    ...  
    return sum;  
}
```

Higher Address

%ebp+12

%ebp+8

%ebp+4


%ebp

%ebp-4

%ebp-8

%ebp-12

Lower Address



Frame pointer %ebp

Stack pointer %esp

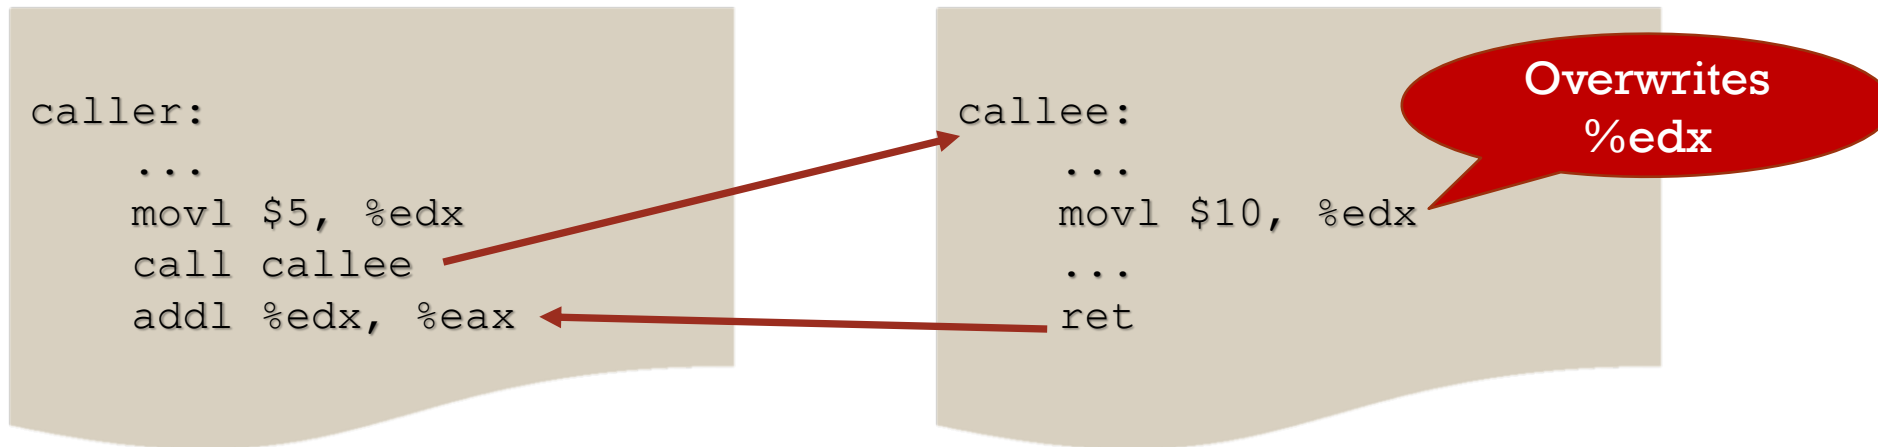
Stack
grows

HOW TO HANDLE RETURN VALUES?

- Most programming languages support only one return value from a function
- Most of the implementations hence support return value using a register
- x86 uses register `%eax` for return value.
- **BUT Wait, what about larger objects say like struct returned by function?**
 - x86 requires caller to allocate space and pass a pointer to that as a “hidden” parameter on stack. Callee will write return value to this address.

MECHANICS FOR REGISTER USAGE

- Both caller and callee functions are running on same HW sharing HW state and resources.
- So what happens when both caller and callee need to make use of same HW registers?




- Somehow `%edx` must be saved so that caller function can use it after call instruction
- **Whose responsibility to save `%edx`?**
 - Need some coordination between caller and callee

REGISTER SAVING CONVENTIONS

- Requirement:
 - Caller function should save (and restore) registers it is using before and after call; callee function should save (and restore) registers it is using in function body.
- Possible redundancy?
 - If callee function is using a register that is saved and restored by caller, does callee function really need to save (and restore) it?
 - If caller function is using a register saved and restored by callee, does caller function really need to save (and restore) it?
- Does caller know which registers are used by callee? Does callee know which registers are used by caller?
 - Inter-procedural analysis and optimizations?
 - What to do in Separate compilation?

CONVENTION: CALLER SAVES ALL REGISTERS


- Caller function saves all registers before calling anything and restores them after call is done



- Every caller function needs to do save and restore at every invocation of call
- May end up saving and restoring registers that are not used by callee function at all
- **Does caller really need to save and restore “ALL” registers?**
 - What if a register is used by caller only before call but is not used after call (i.e. dead register)?


CONVENTION: CALLEE SAVES ALL REGISTERS

- Callee function saves all registers before at start of the function and restores them before returning



- Code to do save, restore is at one place only, in the body of callee function
- May still end up saving and restoring registers that are not used by caller function at all or registers that are “dead”?
- **Does callee really need to save and restore “ALL” registers?**
 - What if a register is not at all used by callee function?

CONVENTION: TRY TO GET BEST OF BOTH WORLDS



REGISTER SAVING CONVENTIONS

- Split available registers into two sets
- *Caller saved Registers*
 - Will be caller's responsibility to save these registers before calling the function and restoring them after the call
 - Callee function can use and modify these registers freely
 - Also known as “*scratch registers*”
- *Callee saved Registers*
 - Will be callee's responsibility to save these registers before using them and restoring them before returning
 - Caller function expects value of these registers will be after function call
 - Also known as “*preserved registers*”
- Yet another **agreement** between caller and callee functions.

REGISTER SAVING CONVENTIONS FOR X86

- Caller saved:
 - `%eax, %ecx, %edx`
- Callee saved:
 - `%ebx, %esi, %edi`
- Special callee saved:
 - `%esp, %ebp`

CONTRACTS BETWEEN CALLER AND CALLEE

- So far, we've seen following things that caller and callee need to agree upon
 - Where is return address located?
 - x86 stored this on stack as part of call instruction
 - How are parameters passed and in what order?
 - x86 passed them on stack and calling conventions for C passed them from right to left order
 - Where is return value?
 - x86 used `%eax` to return value, with larger structs being handled as hidden parameters
 - Which registers are caller saved and which are callee saved
 - x86 had following conventions
 - Caller saved: `%eax`, `%ecx`, `%edx`
 - Callee saved: `%ebx`, `%esi`, `%edi`
 - Special callee saved: `%esp`, `%ebp`
- Such contracts between caller and callee functions are referred as *"Calling Conventions"*

CALLING CONVENTIONS

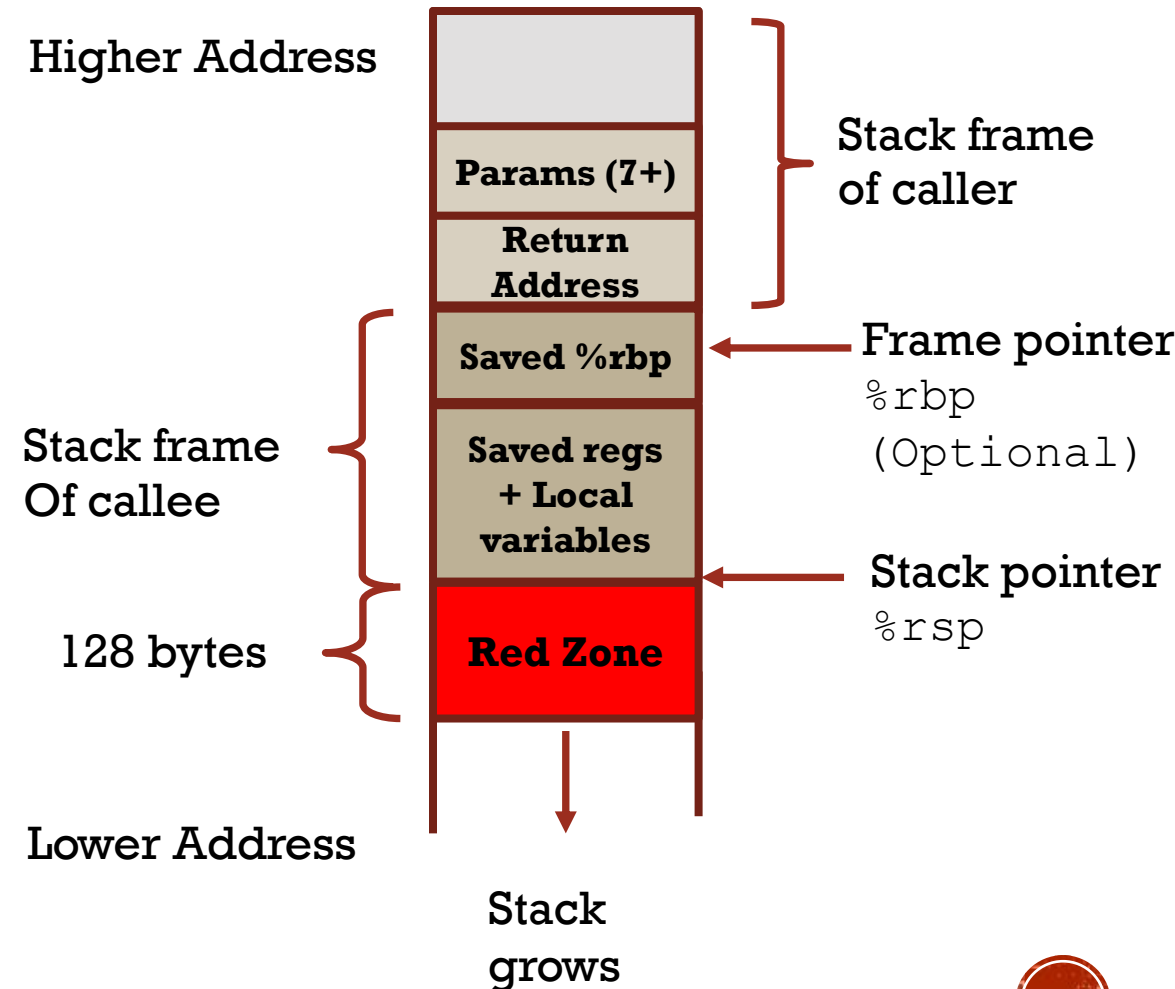
- **Who decides calling conventions?**
 - Application Binary Interface (ABI)
 - Every HW and Operating System specify ABI to be followed for that platform.
 - Contains more than just calling conventions.
 - Compiler implementation for a platform conforms to the ABI of that platform.
- Is there only one calling convention specified by ABI?
 - **What are issues of having only one calling convention?**
 - Platforms support different calling conventions as well mostly for performance reasons e.g. Windows supports `__fastcall` that passes first two arguments in registers `%ecx`, `%edx`, supports bunch of more `stdcall`, `safecall` etc.
- **Can calling conventions be relaxed by compiler as optimization?**
 - What if all callers of a function are known?
 - `static` functions?
 - What if compiler knows it is whole program compilation?
- **What constructs in C language make it hard for compiler to apply optimization related to relaxing calling conventions?**
 - Extern function calls (separate compilation)
 - Function pointers

X86_64 CALLING CONVENTIONS

- x86 had only 8 general purpose registers.
- x86_64 extends them to 64bits (prefix “r” instead of “e”), also adds another 8 general purpose registers r8..r15
- Below details are for System V AMD64 ABI.
 - Linux follows this, Windows follows different ABI
- x86_64 calling convention makes use of more available registers and allows passing up to 6 function parameters in registers and remaining are passed via stack.
 - `%rdi`, `%rsi`, `%rdx`, `%rcx`, `%r8`, `%r9` hold first 6 parameters of a function.
- AMD64 ABI makes frame pointer optional i.e. `%rbp` can be used as general purpose.
 - References are done relative to stack pointer `%rsp`
- `%rax` continues to be return value register

X86_64 STACK FRAME FOR AMD64 ABI

- Remember first 6 parameters are not in stack
- Red Zone:
 - AMD64 ABI defines 128 bytes area beyond stack pointer as “Red Zone”
 - Optimization for leaf functions
 - Leaf functions can use red zone for scratch data without adjusting stack pointer
 - Saves two instructions as don't need to decrement `%rsp` and restore it
 - Red zone will NOT be clobbered by interrupt and signal handlers but will be clobbered by function calls hence useful only in leaf functions



WHAT ELSE FOR FUNCTION CALLS?

- Function Overloading
- Templates
- Virtual Functions
- Exception Handling and stack unwinding

WHAT OTHER TOPICS IN PROG EXECUTION ENV

- Object File formats
- Linker
 - Symbol Resolution
 - Merging object files
- Libraries
 - Static Libraries
 - Shared/Dynamic Libraries
 - Position Independent Code
- Loader
 - Process
 - Virtual Memory
 - C Runtime